

# AI Engineering Project

Final Presentation

Group 4 (Jungo, Luegmayer, Müllner & Schweitzer)

# Project - Cats vs Dogs



# Dataset

- TensorFlow Dataset - (cats\_vs\_dogs)
  - **25.000** Pictures in the dataset
  - **23.262** Valid pictures
  - **1.738** Corrupted images that are dropped



No pre trained Model

# No pre trained Model (5)

Model	Augmentation	BatchNorm	Conv Filters	Dense Units	Params	Train Acc	Val Acc	Overfitting	Notes
Base	No	No	32 → 64 → 64	64	~1.4M	97.12%	84.72%	High	Overfitting visible
Mod-1	No	No	8 → 16 → 32	32	~330k	94.01%	81.62%	Reduced	Simpler, still overfitting
Mod-2	No	No	4 → 8 → 16	16	~84k	75.66%	75.95%	None	Very balanced, efficient
Mod-3	Yes	No	4 → 8 → 16	16	~84k	71.78%	75.19%	None	Augmentation helps generalization
Mod-4	Yes	Yes	8 → 16 → 32	32	~332k	78.97%	79.90%	None	Best generalization and lowest loss



# No pre trained Model - Result

Base  
First Model



Final Training Metrics:  
Training Accuracy: 0.9712  
Validation Accuracy: 0.8472  
Training Loss: 0.0703  
Validation Loss: 0.7046

Modified 4  
Last Model



Final Training Metrics:  
Training Accuracy: 0.7897  
Validation Accuracy: 0.7990  
Training Loss: 0.4518  
Validation Loss: 0.4230

# Data Augmentation

# Data Augmentation

```

Data Augmentation

# Data preprocessing
IMG_SIZE = 160
BATCH_SIZE = 32
AUTOTUNE = tf.data.AUTOTUNE

# Preprocessing (resize + normalize)
def preprocess(image, label):
    image = tf.cast(image, tf.float32)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    image = image / 255.0
    return image, label

# Data-Augmentation pipeline
data_augmentation = tf.keras.Sequential([
    # Randomly crop somewhere between 60-100% of the image, then resize
    tf.keras.layers.RandomCrop(IMG_SIZE, IMG_SIZE),
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomRotation(0.1),
    tf.keras.layers.RandomZoom(0.1, 0.1),
])
```





Pre trained Model

# Pre trained Model

Model	Pretrained	Augmentation	BatchNorm	Conv Filters	Dense Units	Params	Train Acc	Val Acc	Val Loss	Notes
Base	No	No	No	32 → 64 → 64	64	~1.4M	97.12%	84.72%	0.7046	Strong overfitting
Mod-1	No	No	No	8 → 16 → 32	32	~330k	94.01%	81.62%	0.6774	Reduced capacity
Mod-2	No	No	No	4 → 8 → 16	16	~84k	75.66%	75.95%	0.4993	Lightweight, well-regularized
Mod-3	No	Yes	No	4 → 8 → 16	16	~84k	71.78%	75.19%	0.5266	Better generalization
Mod-4	No	Yes	Yes	8 → 16 → 32	32	~332k	78.97%	79.90%	0.4230	Best non-pretrained model
Pretrained	Yes	Yes	—	MobileNetV2	128	~2.4M	97.31%	<b>94.88%</b>	<b>0.3365</b>	Best overall

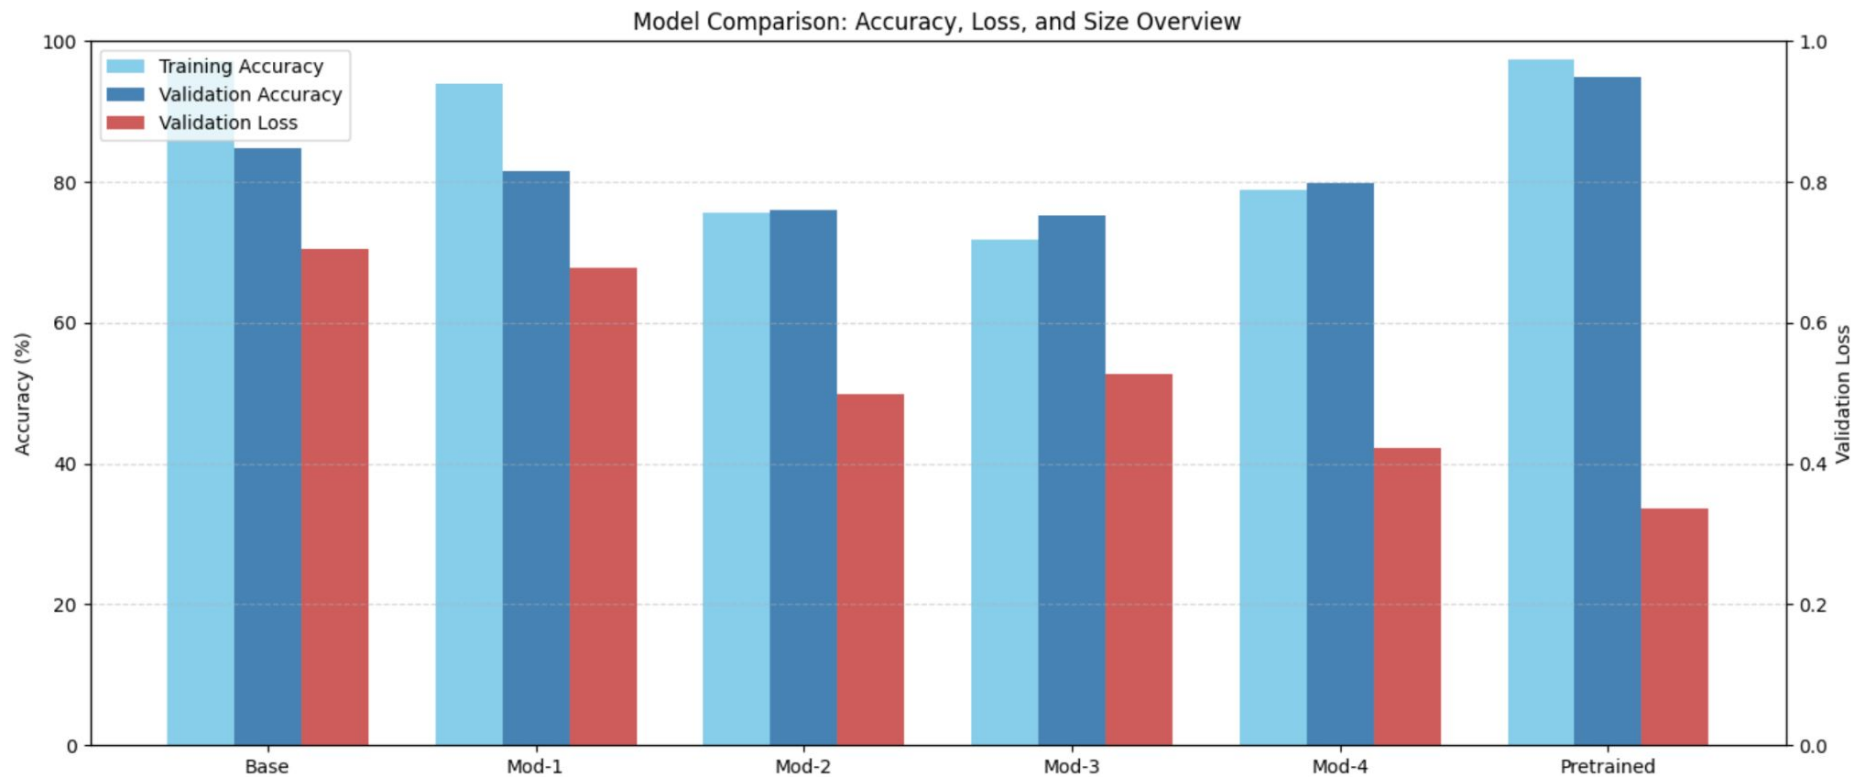
# Pre trained Model




Final Training Metrics:  
Training Accuracy: 0.9731  
Validation Accuracy: 0.9488  
Training Loss: 0.0761  
Validation Loss: 0.3365

# Overall Comparison

# Overall Comparison



# Overall Comparison



# Overall Comparison

